

IBL Computational Neuroscience Course: Introduction to Statistical Model Fitting

Luigi Acerbi

Department of Basic Neuroscience, University of Geneva
Center for Neural Science, New York University
International Brain Lab



April 2020

1 Introduction

- Of models and likelihoods
- The psychometric function

2 Model fitting

- A statistical estimation problem
- Model fitting via optimization
- Optimization algorithms
- Optimization cheat sheets

3 Advanced (Bayesian) topics

- Bayesian Optimization
- Bayesian model fitting

1 Introduction

- Of models and likelihoods
- The psychometric function

2 Model fitting

- A statistical estimation problem
- Model fitting via optimization
- Optimization algorithms
- Optimization cheat sheets

3 Advanced (Bayesian) topics

- Bayesian Optimization
- Bayesian model fitting

What is a model?



The best material model of a cat is another, or preferably the same, cat.

Wiener, *Philosophy of Science* (1945) (with Rosenblueth)

What is a mathematical model?

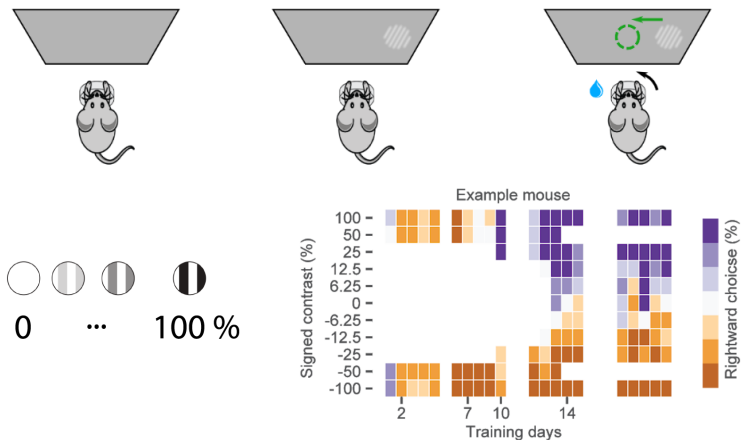
- Quantitative stand-in for a theory
- A *family of probability distributions* over possible datasets:

$$p(\text{data}|\boldsymbol{\theta})$$

- ▶ data is a dataset with n data points (e.g., trials)
 - ▶ $\boldsymbol{\theta}$ is a parameter vector
- **Why?** Description, prediction, and explanation
- Defining $p(\text{data}|\boldsymbol{\theta})$ is the core of model building
 - ▶ Wait, what?
- **How?** Think about the data generation process!

We need some data

IBL Task



(IBL et al., *bioRxiv*, 2020)

Hacking time I

Let's have a look at the data

Type of models

- Descriptive
- Mechanistic
- Process
- Normative
- ...

NB
DT }

ORIGINAL ARTICLE


Commentary

Appreciating the variety of goals in computational neuroscience

Konrad P. Kording PhD¹ | Gunnar Blohm PhD² | Paul
Schrater PhD³ | Kendrick Kay PhD⁴

<https://arxiv.org/abs/2002.03211>

The psychometric function



- Data: (signed contrast, choice) for each trial
- Parameters θ : (μ , σ , γ_{low} , γ_{high})

$$p(\text{rightward choice} | s, \theta) = \gamma_{low} + (1 - \gamma_{low} - \gamma_{high}) \cdot F(s; \mu, \sigma)$$

The psychometric function (alt version)

- Default decision process $F(s; \mu, \sigma)$
- Lapses with probability $\lambda \in [0, 1]$ (*lapse rate*)
- If lapse, respond 'rightward' with probability $\gamma \in [0, 1]$ (*lapse bias*)
- Parameters θ : $(\mu, \sigma, \lambda, \gamma)$


$$p(\text{rightward choice} | s, \theta) = \lambda\gamma + (1 - \lambda) \cdot F(s; \mu, \sigma)$$

Hacking time II

Let's code a psychometric function

Metric for model fitting

We need a quantity to measure *goodness of fit*



- Mean-squared error?
- The likelihood $p(\text{data}|\theta)$

The (log) likelihood

- $p(\text{data}|\theta)$ is a *probability density* as you vary data for a fixed θ
- $p(\text{data}|\theta)$ is the *likelihood*, a function of θ for fixed data

- For numerical reasons we work with $\log p(\text{data}|\theta)$
- Using the rules of probability and logarithms:

$$\begin{aligned}\log p(\text{data}|\theta) &= \log p(\mathbf{r}^{(1)}, \dots, \mathbf{r}^{(n)} | \mathbf{s}^{(1)}, \dots, \mathbf{s}^{(n)}, \theta) \\ &= \log \prod_{i=1}^n p_i \left(\mathbf{r}^{(i)} | \mathbf{r}^{(1)}, \dots, \mathbf{r}^{(i-1)}, \mathbf{s}^{(1)}, \dots, \mathbf{s}^{(n)}, \theta \right) \\ &= \sum_{i=1}^n \log p_i \left(\mathbf{r}^{(i)} | \mathbf{r}^{(1)}, \dots, \mathbf{r}^{(i-1)}, \mathbf{s}^{(1)}, \dots, \mathbf{s}^{(n)}, \theta \right)\end{aligned}$$

- Simplest case: $\log p(\text{data}|\theta) = \sum_{i=1}^n \log p_i(\mathbf{r}^{(i)} | \mathbf{s}^{(i)}, \theta)$
- Model building: Write function with
 - ▶ Input: θ and data
 - ▶ Output: $\log p(\text{data}|\theta)$

Hacking time III

Let's code up a log-likelihood function

1 Introduction

- Of models and likelihoods
- The psychometric function

2 Model fitting

- A statistical estimation problem
- Model fitting via optimization
- Optimization algorithms
- Optimization cheat sheets

3 Advanced (Bayesian) topics

- Bayesian Optimization
- Bayesian model fitting

Model fitting

Model fitting \sim *statistical estimation* problem

1. Maximum likelihood estimation (MLE)

- Find maximum of $p(\text{data}|\theta)$

$$\hat{\theta}_{\text{ML}} = \arg \max_{\theta} p(\text{data}|\theta) = \arg \max_{\theta} \log p(\text{data}|\theta)$$

2. Bayesian posterior

$$p(\theta|\text{data}) = \frac{p(\text{data}|\theta)p(\theta)}{p(\text{data})} \propto p(\text{data}|\theta)p(\theta)$$

- For $n \rightarrow \infty$ converges to MLE (if $p(\hat{\theta}_{\text{ML}}) \neq 0$)
- Full posterior*: informative about parameter uncertainty and trade-offs
- Maximum-a-posteriori (MAP)*: $\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} p(\theta|\text{data})$

How to do model fitting?

Maximum likelihood estimation (MLE), Maximum-a-posteriori (MAP)

- Model fitting \sim *optimization problem*


Bayesian posterior

- How do we represent/approximate an arbitrary posterior distribution?
 - ① Use a known (easier) distribution (*variational inference*)
 - ② Use a bunch of discrete samples (*Markov-Chain Monte Carlo*)

Model fitting via optimization


- Find single θ that best describes the data
- (For this section we switch notation from θ to \mathbf{x})
- Given $\tilde{f}(\mathbf{x}) \equiv \begin{cases} \log p(\text{data}|\mathbf{x}) & \text{maximum likelihood} \\ \log p(\text{data}|\mathbf{x}) + \log p(\mathbf{x}) & \text{maximum-a-posteriori} \end{cases}$
- By convention, we *minimize* $f(\mathbf{x}) \equiv -\tilde{f}(\mathbf{x})$
- \implies Find $\mathbf{x}_{opt} \approx \arg \min_{\mathbf{x}} f(\mathbf{x})$ as fast as possible
- General case: $f(\mathbf{x})$ is a *black box*
 - ▶ Sometimes we can compute the gradient

How hard can it be?









Source: Wikimedia Commons

How hard can it be?









Source: Wikimedia Commons

How hard can it be?



How hard can it be?



How hard can it be?

neval	x_1	x_2	$f(x)$
1	-0.500	2.500	508.500
2	-0.525	2.500	497.110
3	-0.500	2.625	566.313
4	-0.525	2.375	443.063
5	-0.537	2.250	386.953
6	-0.563	2.250	376.320
7	-0.594	2.125	316.702
8	-0.606	1.875	229.824
9	-0.647	1.563	133.598
10	-0.703	1.438	91.847
11	-0.786	1.031	20.292
12	-0.839	0.469	8.918
13	-0.962	-0.359	168.785
14	-0.978	-0.063	107.796
15	-0.895	0.344	24.553
16	-0.730	1.156	41.905
17	-0.854	0.547	6.760
18	-0.907	-0.016	73.917
19	-0.816	0.770	4.366
20	-0.831	0.848	5.818
21	-0.793	1.070	22.655
22	-0.839	0.678	3.448
23	-0.824	0.600	3.955
24	-0.846	0.508	7.766
25	-0.824	0.704	3.391
26	-0.839	0.782	4.004
27	-0.828	0.645	3.497
28	-0.835	0.737	3.523
29	?	?	?

Optimization can be hard

- ① Optimizer does not see the landscape!
- ② Multiple local minima or saddle points ('non-convex')
- ③ Expensive function evaluation
- ④ Noisy function evaluation
- ⑤ Rough landscape (numerical approximations, etc.)

Optimization algorithms

Gradient-based methods

- Stochastic gradient descent (e.g., ADAM)
- Quasi-Newton methods (e.g., BFGS aka `fminunc`/`fmincon`)


Gradient-free methods

- Nelder-Mead (`fminsearch`)
- Pattern/direct search (`patternsearch`)
- Simulated annealing
- Genetic algorithms
- CMA-ES
- Bayesian optimization
- Bayesian Adaptive Direct Search (BADS; Acerbi & Ma, *NeurIPS* 2017)

Demos: <https://github.com/lacerbi/optimviz>

Nelder-Mead (`fminsearch`)


J. A. Nelder & R. Mead, A simplex method for function minimization (1965)



Source: Encyclopedia of Artificial Intelligence (2009)


Bounded optimization: `fminsearchbnd` (John d'Errico)

Newton method



Source: StackExchange


Newton method



Source: StackExchange

Needs the inverse of the curvature (inverse Hessian)
Very expensive in high dimension

Quasi-Newton methods (fminunc, fmincon)




Source: StackExchange

Approximate Hessian (DFP) or inverse Hessian (BFGS) via gradient
Very fast and efficient on smooth problems

Direct search (patternsearch)

R. Hooke and T.A. Jeeves, “Direct search” solution of numerical and statistical problems (1961)




Source: Wikimedia Commons

Genetic Algorithms (ga)

J.H. Holland, *Adaptation in Natural and Artificial Systems* (1975)

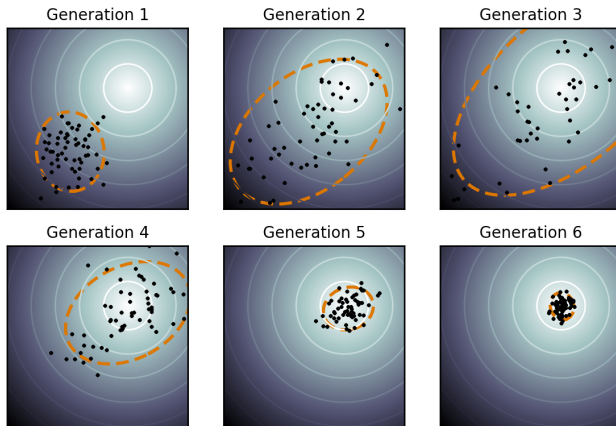
- Evolutionary algorithm
- Population based



Source: An Educational GA Learning Tool (IEEE)

Cov. Matrix Adaptation - Evolution Strategies (cmaes)

[*] N. Hansen, S. D. Müller, P. Koumoutsakos, Reducing the time complexity of the derandomized evolution strategy with covariance matrix adaptation (CMA-ES), (2003)



Hacking time IV

Let's optimize the log-likelihood for the psychometric model

Optimization cheat sheet, page 1

Rule zero


Understand your problem \implies often a *gray box*

Input variables:

- Dimensionality: low ($D \lesssim 10$) or high ($D \gg 20$)
- Bounds: Think of *hard* and *plausible* bounds
- Parameterization: Not all parameterizations are created equal

Target function:

- Convexity: convex or non-convex
- Smoothness: smooth or rough



- Deterministic or stochastic
 - ▶ If stochastic \implies minimize $\mathbb{E}[f(\mathbf{x})]$
- Computational cost:
cheap ($\ll 0.01$ s), moderate (0.01-1 s), or expensive ($\gg 1$ s)

Fundamental theorem

'No Free Lunch' theorem \implies no single best optimizer for all problems
(But not all methods are created equal!)

- Is your problem smooth?
 - ▶ If you have the gradient \implies BFGS
 - ▶ If low- D and cheap \implies BFGS with finite differences
 - ▶ If low- D and (moderately) costly \implies BADS
- Is your problem rough or noisy?
 - ▶ First, try and make it smooth and deterministic!
 - ▶ If gradient is available and high- D \implies SGD (e.g., ADAM)
 - ▶ If high- D and cheap \implies CMA-ES
 - ▶ If low- D and (moderately) costly \implies BADS
- Is your problem high- D , costly, and you do not have the gradient?
 - ▶ Give up and pray

Optimization cheat sheet, page 3

The golden rule

No optimizer can guarantee to find the global optimum


⇒ **Always** perform multiple distinct optimization runs ('restarts')

- How to choose starting points?

- ▶ Draw from prior distribution
- ▶ Draw from a 'plausible' box
- ▶ Sieve method
- ▶ Use space-filling designs (quasi-random sequences)



Random



Space-filling

- How many restarts?

- ▶ As many as you need
- ▶ Informally, check that 'most' points converge to the same solution
- ▶ *Bootstrap* approach (Acerbi, Dokka et al., *PLoS Comp Biol* 2018)

1 Introduction

- Of models and likelihoods
- The psychometric function

2 Model fitting

- A statistical estimation problem
- Model fitting via optimization
- Optimization algorithms
- Optimization cheat sheets

3 Advanced (Bayesian) topics

- Bayesian Optimization
- Bayesian model fitting

Bayesian Optimization

- 1 Start with a prior over functions (Gaussian process)
- 2 Find $\tilde{\mathbf{x}}$ that maximizes *acquisition function* (exploration/exploitation)
- 3 Evaluate $f(\tilde{\mathbf{x}})$
- 4 Compute posterior over functions (Gaussian process)
- 5 goto 2

J. Mockus, *Journal of Global Optimization* (1994)

Review paper: Frazier (2018) <https://arxiv.org/abs/1807.02811>

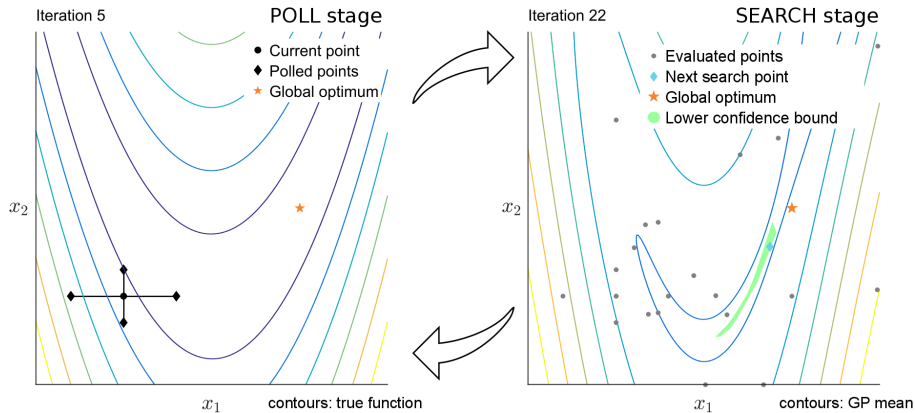
Why don't we use Bayesian optimization *all the time*?

- Matrix inversion is $O(n^3)$
- Model mismatch



from xkcd.com/2048

Bayesian Adaptive Direct Search (bads)



Acerbi & Ma, *NeurIPS* (2017)

BADS summary

- Good for moderately costly ($\gtrsim 0.1$ s) or noisy functions
- Scales okay with n (uses only local neighborhood)
- Local approximation deals with nonstationarity
- Explicit support for noise
- Outperforms other algorithms (Acerbi & Ma, 2017)

The screenshot shows the GitHub repository page for 'lacerbi / bads'. At the top, there are buttons for 'Unwatch' (13), 'Star' (135), and 'Fork' (25). Below these are navigation tabs: 'Code', 'Issues' (4), 'Pull requests' (0), 'Actions', 'Projects' (0), 'Wiki', 'Security', 'Insights', and 'Settings'. The repository name 'Bayesian Adaptive Direct Search (BADS) optimization algorithm for model fitting in MATLAB' is displayed, with an 'Edit' button. Below the name are topic tags: 'optimization-algorithms', 'bayesian-optimization', 'log-likelihood', 'noiseless-functions', 'noisy-functions', 'matlab', and 'Manage topics'. At the bottom, statistics are shown: 161 commits, 2 branches, 0 packages, 6 releases, 1 contributor, and GPL-3.0 license.

lacerbi / bads

Unwatch 13 Star 135 Fork 25

<> Code Issues 4 Pull requests 0 Actions Projects 0 Wiki Security Insights Settings


Bayesian Adaptive Direct Search (BADS) optimization algorithm for model fitting in MATLAB Edit

optimization-algorithms bayesian-optimization log-likelihood noiseless-functions noisy-functions matlab Manage topics

161 commits 2 branches 0 packages 6 releases 1 contributor GPL-3.0


<https://github.com/lacerbi/bads>

Bayesian posteriors



$n = 1353$ trials

Bayesian posteriors



$n = 90$ trials

Benefits of Bayesian posteriors

- Check for parameter uncertainty, trade-offs, identifiability
 - ▶ Deeper understanding of your model
 - ▶ Robustness of claims (Acerbi, Ma, Vijayakumar, *NeurIPS* 2014)
- Less overfitting
- Use posterior samples to compute model comparison metrics
 - ▶ DIC, WAIC, LOO-CV
- Fully taking into account uncertainty is just *better*

How do I get Bayesian posteriors?

- Markov Chain Monte Carlo (MCMC; e.g. slice sampling, NUTS)
- Variational inference

Applied example

 OPEN ACCESS  PEER-REVIEWED

RESEARCH ARTICLE

Bayesian comparison of explicit and implicit causal inference strategies in multisensory heading perception

Luigi Acerbi  , Kalpana Dokka , Dora E. Angelaki, Wei Ji Ma

Published: July 27, 2018 • <https://doi.org/10.1371/journal.pcbi.1006110>

Final slide

- Contact me at luigi.acerbi@internationalbrainlab.org
- Python tutorial: github.com/lacerbi/ibl-2020-tutorial
- Optimization demos: github.com/lacerbi/optimviz

MATLAB toolboxes:

- BADS available at github.com/lacerbi/bads
- VBMC available at github.com/lacerbi/vbmc

Thanks!

(Time for questions?)